



Obsidian Is My AI's Second Brain. Here's My Full Setup.



I tested every obsidian + claude setup - here's the winner
AI with Remy



Titta på



ps; give this guide to claude code/ cowork and ask it to explain anything you're confused about, and to set this up for you :)

make sure you understand tho, what it's doing, don't just blindly set it up.

Every morning, Claude plans my day. Automatically. I wake up, open my notes, and there's a briefing already written.

Last week it said: *"You've been talking about rewriting your welcome email since January. You've mentioned it in four different notes. You haven't started. That's your priority today."*

Or the other day I was building a landing page and Claude went into my vault, found a transcript from a client call three months ago, and pulled out a sentence to use as a testimonial on the page. I'd completely forgotten that call even happened.

That's what happens when your AI actually has context.

The tool behind all of this is Obsidian. It's a stupid simple note-taking app. But when you plug it into Claude Code, it becomes something else entirely. An entire neural network for your AI agent. It connects all your thoughts (projects, people, meetings, brand docs, strategies) in a web of linked markdown files.

And Claude can traverse all of it. I connect this same vault to Claude Code, Hermes agents, and OpenClaw. Same markdown files, same context, different harnesses.

This is my full setup.

The problem: every session is groundhog day

Most people don't think about this when they start using AI agents. Every time you start a new session, your AI remembers nothing. Your brand voice, your customer profile, your meeting notes, your projects. Gone. You're re-explaining yourself constantly.

Context is the real bottleneck, not the models.

The models are smart enough. They've been smart enough for a while. What they're missing is *you*. Your business, your decisions, your history, your thinking. Without that, you're just talking to a very capable stranger every single time.

I tried solving this a few different ways. I tried cramming everything into CLAUDE.md files. I tried keeping context docs in every project folder. I even tried pointing agents at Notion.

The Notion approach was the closest. Keep everything in one central place and just point each agent there. But that meant every time your agent needed to understand something about you, which is every single chat, it had to hit the Notion API to go find it. Slow & Token-heavy. And it wouldn't be proactive finding new information (like my meeting notes for the testimonial example from earlier).

Then I found Obsidian.

What is Obsidian

If you've never used it, Obsidian is a note-taking app. Like Apple Notes or Notion.

The key difference: every note is just a markdown file sitting on your computer. There's no database, no proprietary format, no cloud lock-in. Just folders and  files.

You open Obsidian and it points at a folder on your machine. That folder is called your "vault." Every note you create is a plain text file inside that folder. You can open them in any text editor. You can back them up however you want. You can move them, rename them, organise them with regular folders.

If Obsidian disappeared tomorrow, your notes would still be right there on your computer. Try saying that about Notion.

Why Obsidian over everything else

I've used Notion, Apple Notes, Google Docs, Roam. None of them come close for what I'm doing. And it all comes down to one thing: Obsidian is the best note-taking app for AI.

A few reasons:

It's all markdown. AI reads markdown natively. It's the language AI agents already think in. [CLAUDE.md](#), [AGENTS.md](#), memory files, skills. Everything in the AI world is markdown. Your notes should be too.

It's local files. Your AI agent can traverse your file system lightning fast. There are no API calls, no network latency, no rate limits. Claude Code can read a note from your vault in milliseconds. Compare that to waiting for a Notion API response.

One source of truth. I keep all my brand docs, customer profiles, meeting notes, and strategies in one vault. Every project and every agent points there. When I update my brand voice doc, every agent sees the change immediately. You're not syncing anything or duplicating files across 20 different folders.

Wikilinks connect everything. In Obsidian, you can link any note to any other note with `[[double brackets]]`. Your meeting notes link to your project notes. Your project notes link to your brand docs. It creates a web of connected knowledge that your AI can follow.

It's free. And your data is yours. Forever.

My vault structure

This is the actual folder structure I use:

```
Vault/ |— Areas/ Business areas & clients | |— AI with Remy/ | |
|— _context/ Brand voice, ICP, positioning, etc. | | |— Content/
All content, organised by type | | |— newsletters/ | | |—
youtube/ | | |— articles/ | | |— reels/ | | |— Notes/ Junk
drawer (research, references, anything) | | |— Projects/ Per-area
project tracker | |— Client Name/ One folder per client | |—
_context/ Client business docs | |— Notes/ | |— Projects/ | |—
Daily logs/ Daily notes + Backlog (pinned at top) |— Meeting Notes/
Auto-synced from Granola |— Notes/ Evergreen reference notes | |—
_System/ |— Templates/ Auto-applied templates |— attachments/
```

Let me break down the key pieces.

Areas

Each business or client gets its own folder under `Areas/`. This is where all the context about that business lives. I run multiple brands and take on consulting clients, so having clean separation here matters a lot.

The `_context/` pattern

This is the most important part of my setup.

Every area has a `_context/` folder that contains the canonical docs about that business. For AI with Remy, that's:

- Brand Voice Guide
- Brand Guidelines (colours, typography, design system)
- Ideal Customer Profile
- Positioning (competitive landscape, market position)
- Product & Service Information
- Newsletter Structure Guide
- Manifesto

These are the single source of truth for everything about the business. They never get duplicated. Every AI agent, every workspace, every project just points here.

When something changes, my pricing, my positioning, my ICP, I update one file in `_context/` and every agent that references it gets the update automatically. No hunting through 20 different folders to find and update the same doc.

Projects

Each area has its own `Projects/` folder. One markdown file per project with a single property: `status` (`active` , `paused` , `done`). An Obsidian Bases kanban in each area shows me what's active at a glance. No priority field, no tags, no due dates — if something matters, work on it.

Daily logs

One note per day, auto-created from a template. I dump thoughts, tasks, and notes throughout the day. These are incredibly valuable for the AI because they capture what I'm actually thinking about and working on.

Meeting Notes

These auto-sync from Granola, a meeting transcription app. Every call gets transcribed and dropped into this folder. This is how Claude found that client testimonial I'd forgotten about. It was sitting in a meeting transcript from three months ago.

Notes

Evergreen reference notes. SOPs, setup docs, research, ideas. Anything that's not tied to a specific date or project.



One thing worth saying: this is my exact setup. I keep content planning, project management, and client work all inside Obsidian because I like having everything in one place. But there's no right or wrong way to do this. If you prefer tracking projects in Notion or ClickUp or Linear, keep doing that. The part that matters for the AI connection is the `_context/` docs, the daily logs and meeting notes.

The rest is personal preference.

What day-to-day Obsidian use looks like

I'm not spending hours organising notes. The whole point is that this is low friction.

Morning: I have a Claude skill that runs on a schedule every morning. It checks my calendar, triages my emails, reviews my recent daily logs and project notes, then writes a full morning briefing into today's daily note. By the time I open Obsidian, there's already a summary of my schedule, priorities, and tasks waiting for me. I add to it throughout the day.

Meetings: Granola runs in the background, transcribes my calls, and auto-syncs the transcripts into `Meeting Notes/`. I don't touch anything.

Working on a project: Notes, research, and ideas go into the relevant Area folder. If I'm brainstorming a newsletter topic, it goes in `Content/newsletters/`. If I'm working on something that doesn't fit Content or Projects (research, product docs, references), it goes in the area's `Notes/` folder.

Content tracking: Every content note has just two frontmatter properties: `type` (newsletter/youtube/article/reel) and `status` (idea/draft/published). An Obsidian Bases kanban shows me what's in progress, what's published, what's just an idea. No tags, no dates in frontmatter — the filename and folder do that work.

Updating business docs: When something changes in the business (new pricing, refined positioning, updated ICP) I edit the file in `_context/` directly. Done. Every agent sees it.

The theme here is: I write notes for myself, and the AI benefits as a side effect. I'm not maintaining a database for my AI. I'm taking notes like a normal person, and the structure makes it useful for both humans and machines.

Connecting Obsidian to Claude Code

This is probably the bit you actually care about. There are three layers to how Claude Code accesses my vault, from simplest to most advanced.

Layer 1: `CLAUDE.md` imports (passive context)

Claude Code has a config file called CLAUDE.md that it reads automatically at the start of every session. There are two levels: a global one that applies everywhere, and project-level ones for specific workspaces.

The global CLAUDE.md lives at `~/.claude/CLAUDE.md`. This is the most important file in my entire setup. It tells Claude how I work, where my vault is, how to search it, and sets the rules for every single session regardless of which project I'm in. This is where I point Claude at my Obsidian vault and teach it to use the CLI, the scheduled tasks, how I like to work. Anything that should be true across all projects goes here.

This is the exact Obsidian section from mine. You can copy this and adapt it to your vault:

```
## Obsidian – Second Brain Obsidian is the single source of truth for all project context, brand information, strategy docs, meeting notes, and business knowledge.
```

I also have **global rules** in `~/.claude/rules/`. These are separate markdown files for specific topics that Claude loads alongside the global CLAUDE.md. I have one for Obsidian (when to use the CLI vs semantic search, how to load context), one for Telegram (how to handle messages from my bot), one for Google Workspace (which accounts to use). They keep the global CLAUDE.md from becoming a giant mess.

My Obsidian rules file at `~/.claude/rules/obsidian.md` looks like this:

```
# Obsidian Rules ## Context Loading 1. Before starting any project work, check Obsidian for relevant context. Don't ask the user for information that's already documented – go find it. 2. Check if the project's CLAUDE.md already has @path imports from the vault – if so, that context is already loaded. 3. If not, use `obsidian search query=<project name>` to find relevant vault notes. 4. Read key context files (brand guidelines, PRDs, strategy docs) before making decisions. 5. Never assume you know the current state of a project – Obsidian has the latest. 6. When the user references a meeting, decision, note, or any business context – search Obsidian first rather than asking the user to explain it. ## CLI vs QMD MCP Routing – Default to CLI for all vault access: reading, writing, targeted keyword searches, daily notes, tasks, tags, backlinks, properties. – Escalate to QMD MCP only when CLI search returns poor/no results, or when searching by concept rather than exact words. – Never use both for the same query – start with CLI, escalate to QMD only if CLI misses.
```

The rules file is where the real behaviour lives. They tell Claude to go search your vault *before* asking you to explain something. That's what makes it feel proactive rather than helpless.

Project-level CLAUDE.md files are where things get specific. Each project gets its own, and the magic is the `@` import syntax. You can point to files in your vault, and Claude loads them as context.

In practice it looks like this:

```
# My Project CLAUDE.md ## Shared Context (auto-loaded from Obsidian vault)
@~/Desktop/Vault/Areas/AI with Remy/_context/Brand Voice Guide.md
@~/Desktop/Vault/Areas/AI with Remy/_context/Ideal Customer Profile.md
@~/Desktop/Vault/Areas/AI with Remy/_context/Positioning.md
```

Three lines. Now every time Claude starts a session in this project, it automatically knows my brand voice, who my customers are, and how I'm positioned in the market. You never have to copy-paste context or re-explain yourself again.

For a client project, it looks like this:

```
## Shared Context (auto-loaded from Obsidian vault)
@~/Desktop/Vault/Areas/Client Name/_context/About.md
@~/Desktop/Vault/Areas/Client Name/_context/Ideal Customer Profiles.md
```

Same pattern, different vault path. The `_context/` folder structure makes this stupid easy. Every area's context is in the same predictable place.

So when Claude starts a session, it loads the global CLAUDE.md (which teaches it how to work with me), the global rules (specific behavioural instructions), and the project CLAUDE.md (which pulls in the right business context from the vault). All of that happens before I type a single word.

The key insight: The context docs live in the vault, not in the project. Twenty different projects can all point to the same brand voice doc. Update it once, update it everywhere.

Layer 2: Obsidian CLI (active search)

Layer 1 gives Claude context that's pre-loaded at the start of every session. But sometimes Claude needs to go *find* something on the fly. A meeting note you forgot about, a project status, something you wrote last Tuesday. That's where the Obsidian CLI comes in.

Obsidian has a built-in command-line interface that lets you read, search, and write to your vault from the terminal. Since Claude Code can run terminal commands, it can use the CLI to search your vault whenever it needs to. You don't have to pre-load everything. Claude just goes and looks.

This is how it found that four-note pattern about the welcome email. It searched my vault, found multiple notes mentioning it across different dates, connected the dots, and surfaced it as a priority.

The CLI commands are already included in the global [CLAUDE.md](#) template above, so Claude knows how to use them out of the box. Full docs here: <https://help.obsidian.md/cli>

Layer 3: QMD MCP (semantic search, optional power-up)

The CLI does keyword search. It's fast and it covers most of what you need. But keywords have limits. If you search for "pricing decision" but the note is titled "Revenue Strategy Meeting," the CLI won't find it.

That's the gap QMD fills. QMD is a separate tool (by Tobi Lutke, the Shopify founder) that adds semantic vector search to your vault. It understands *meaning*, not just words. You install it as an MCP server, which means Claude can call it directly alongside the CLI.

So the two tools serve different purposes:

- **CLI** for targeted lookups where you roughly know what you're looking for (reading notes, keyword search, daily logs, tasks)
- **QMD** for fuzzy conceptual searches where you know a note exists but can't remember the exact words

I start most sessions with just the CLI. QMD kicks in when Claude searches for something and comes back empty, or when I ask a question like "what did we decide about the onboarding flow?" where the answer is spread across multiple notes.

Install QMD:

```
npm install -g @tobylu/qmd qmd collection add ~/Desktop/Vault --name obsidian qmd embed # one-time, builds the search index
```

MCP config (add to `~/.claude.json`):

```
{ "mcpServers": { "qmd": { "command": "qmd", "args": ["mcp"], "env": { "HOME": "/Users/yourname" } } } }
```

Important: keeping QMD up to date

There's one gotcha with QMD that trips people up. When you run `qmd embed`, it takes a snapshot of every note in your vault and converts it into a numerical vector (an "embedding") that captures its meaning. That's how semantic search works — it compares the meaning of your question against the meaning of every note.

But here's the thing: embeddings don't auto-update. If you add or edit a note after the last embed, QMD doesn't know it exists. Your semantic search is only as fresh as your last `qmd embed` run. The keyword search (`qmd search`) also relies on the index being current.

If you're using Obsidian actively, you need to re-embed regularly. The easiest way is to have Claude set up a scheduled task that runs it automatically. Here's a prompt you can paste into Claude Code:

👉 *Set up a macOS Launch Agent called `com.qmd.auto-embed` that runs `qmd update` & `qmd embed` twice daily at 12am and 12pm. Make sure the `PATH` includes the directory where `qmd` is installed (check with `which qmd`). Log output to `/tmp/qmd-auto-embed.log`. Load it with `launchctl` after creating it.*

That gives you a background process that keeps your search index fresh without you ever thinking about it. The embed step is incremental — it only processes new or changed files, so it's fast after the first run.

Storage-wise, QMD uses about 1 GB total. Almost all of that is three small AI models it downloads once (for embeddings, reranking, and query expansion). They run locally on your Mac's GPU — no API calls, no cloud. Your actual vault index is tiny: my 126 notes take up about 12 MB. Even a vault with thousands of notes would stay small.

Honestly, start without QMD. The CLI handles 90% of what you need. Add it when your vault grows large enough that keyword search starts missing things.

Plugins I use

Obsidian has hundreds of community plugins. I keep it minimal:

Plugin	What it does
Templater	Auto-applies templates when I create notes in specific folders. New daily log? Template applied. New content note? Frontmatter pre-filled.
Excalidraw	Whiteboarding and diagrams inside Obsidian. I use it for architecture diagrams and visual planning.
Granola Sync Plus	Auto-syncs meeting transcripts from Granola into my vault. Every call I take shows up in Meeting Notes/ automatically.
Terminal	Open a terminal from within Obsidian. Handy for quick CLI commands without switching apps.

Plus the built-in **Daily Notes** and **Bases** (Notion-style database views over your markdown files).

That's it. Five plugins. The power is in the structure and the AI connection, not the plugins.

What this actually looks like in practice

Let me bring it full circle.

When Claude plans my morning, it's not just reading a to-do list. It's searching my daily logs for patterns. It's checking my project statuses. It's reading recent meeting notes to see what commitments I made. It's looking at my content calendar to see what's due. And then it pulls all of that together into a plan, with the kind of context that makes it feel like it actually knows me.

When I'm building a landing page for a client, Claude loads the client's brand docs from `_context/` automatically. But then it goes further. Searching meeting transcripts for quotes, checking project notes for feature requirements, reading previous content for messaging that's already been validated.

When I'm writing a newsletter, Claude has my brand voice guide, my newsletter structure, and my ICP all loaded from the vault. It knows how I write, who I write for, and what format to follow. I don't need to explain any of that.

This is what people mean when they talk about AI agents with memory. It's just a well-structured folder of markdown files that your AI can read. No fancy feature built into the model.

The harnesses will keep changing. Claude Code, OpenClaw, Cowork, whatever comes next. But your vault, your markdown files, your context docs, your thinking, that's yours. It works with any tool that can read a file. And that's the whole point.

Build your vault. Point your AI at it. Let it remember everything you forget.